

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/> (<https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: 1 %matplotlib inline
2 import warnings
3 warnings.filterwarnings("ignore")
4
5
6 import sqlite3
7 import pandas as pd
8 import numpy as np
9 import nltk
10 import string
11 import matplotlib.pyplot as plt
12 import seaborn as sns
13 from sklearn.feature_extraction.text import TfidfTransformer
14 from sklearn.feature_extraction.text import TfidfVectorizer
15
16 from sklearn.feature_extraction.text import CountVectorizer
17 from sklearn.metrics import confusion_matrix
18 from sklearn import metrics
19 from sklearn.metrics import roc_curve, auc
20 from nltk.stem.porter import PorterStemmer
21
22 import re
23 # Tutorial about Python regular expressions: https://pymotw.com/2/re/
24 import string
25 from nltk.corpus import stopwords
26 from nltk.stem import PorterStemmer
27 from nltk.stem.wordnet import WordNetLemmatizer
28
29 from gensim.models import Word2Vec
30 from gensim.models import KeyedVectors
31 import pickle
32
33 from tqdm import tqdm
34 import os
```

```
E:\anaconda\lib\site-packages\gensim\utils.py:1197: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")
```

```
In [2]: 1 #from google.colab import drive
2 #drive.mount('/content/gdrive')
```


	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all	Th con th arc



```
In [4]: 1 display = pd.read_sql_query("""
2 SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
3 FROM Reviews
4 GROUP BY UserId
5 HAVING COUNT(*)>1
6 """, con)
```

```
In [5]: 1 print(display.shape)
2 display.head()
```

(80668, 7)

Out[5]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

```
In [6]: 1 display[display['UserId']=='AZY10LLTJ71NX']
```

Out[6]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	5

```
In [7]: 1 display['COUNT(*)'].sum()
```

```
Out[7]: 393063
```

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [8]: 1 display= pd.read_sql_query("""
2 SELECT *
3 FROM Reviews
4 WHERE Score != 3 AND UserId="AR5J8UI46CURR"
5 ORDER BY ProductID
6 """, con)
7 display.head()
```

Out[8]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [9]: 1 #Sorting data according to ProductId in ascending order
        2 sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort',
```

```
In [10]: 1 #Deduplication of entries
        2 final=sorted_data.drop_duplicates(subset={"UserId", "ProfileName", "Time", "Text"}, keep='first', inplace=False)
        3 final.shape
```

Out[10]: (364173, 10)

```
In [11]: 1 #Checking to see how much % of data still remains
        2 (final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[11]: 69.25890143662969

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations


```
In [12]: 1 display= pd.read_sql_query("""
2 SELECT *
3 FROM Reviews
4 WHERE Score != 3 AND Id=44737 OR Id=64422
5 ORDER BY ProductID
6 """, con)
7
8 display.head()
```

Out[12]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1	5	1224892800	Bought This for My Son at College
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2	4	1212883200	Pure cocoa taste with crunchy almonds inside

```
In [13]: 1 final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [14]: 1 #Before starting the next phase of preprocessing Lets see the number of entries Left
2 print(final.shape)
3 final = final.sample(100000
4                       )
5
6 #How many positive and negative reviews are present in our dataset?
7 final['Score'].value_counts()
```

(364171, 10)

```
Out[14]: 1 84425
0 15575
Name: Score, dtype: int64
```

```
In [15]: 1 final = final.sort_values('Time',ascending = True)#sorting by time for time series cross validation
        2 final.shape
```

Out[15]: (100000, 10)

```
In [16]: 1 final.head()
```

Out[16]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	
138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	2	1	940809600	This wh
417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	0	1	944092800	
346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	2	1	944438400	
346116	374422	B00004CI84	A1048CYU0OV4O8	Judy L. Eans	2	2	1	947376000	
346141	374450	B00004CI84	ACJR7EQF9S6FP	Jeremy Robertson	2	3	1	951523200	Bettlejuic



[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [17]: 1 # printing some random reviews
2 sent_0 = final['Text'].values[0]
3 print(sent_0)
4 print("="*50)
5
6 sent_1000 = final['Text'].values[1000]
7 print(sent_1000)
8 print("="*50)
9
10 sent_1500 = final['Text'].values[1500]
11 print(sent_1500)
12 print("="*50)
13
14 sent_4900 = final['Text'].values[4900]
15 print(sent_4900)
16 print("="*50)
```

I can remember seeing the show when it aired on television years ago, when I was a child. My sister later bought me the LP (which I have to this day, I'm thirty something). I used this series of books & songs when I did my student teaching for preschoolers & turned the whole school on to it. I am now purchasing it on CD, along with the books for my children 5 & 2. The tradition lives on!

=====

Frank's Red Hot sauce is the best hot sauce there is. It's good on just about anything - even Chinese food!

Seriously, try it on everything! It gives anything a really great flavor & kick.

=====

I purchased these on the special Friday sale where they were \$7.99 for 3 cans. At the store where I normally shop here one can is almost \$5 so this was a fantastic deal. I liked how they were lightly salted and not as salty as a lot of nuts are.

I agree with the other review that they were mostly almonds. I'd say 60% almonds, 30% cashews and 10% macadamias but they were still quite good.

=====

This is the second time that I have ordered Belly Flops. They contain an excellent assortment of flavors. The odd shapes make them interesting, and the lack of the words "Jelly Belly" does not detract at all from the taste!

=====

```
In [18]: 1 # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
2 sent_0 = re.sub(r"http\S+", "", sent_0)
3 sent_1000 = re.sub(r"http\S+", "", sent_1000)
4 sent_150 = re.sub(r"http\S+", "", sent_1500)
5 sent_4900 = re.sub(r"http\S+", "", sent_4900)
6
7 print(sent_0)
```

I can remember seeing the show when it aired on television years ago, when I was a child. My sister later bought me the LP (which I have to this day, I'm thirty something). I used this series of books & songs when I did my student teaching for preschoolers & turned the whole school on to it. I am now purchasing it on CD, along with the books for my children 5 & 2. The tradition lives on!

```
In [19]: 1 # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-tags-from-an-element
2 from bs4 import BeautifulSoup
3
4 soup = BeautifulSoup(sent_0, 'lxml')
5 text = soup.get_text()
6 print(text)
7 print("="*50)
8
9 soup = BeautifulSoup(sent_1000, 'lxml')
10 text = soup.get_text()
11 print(text)
12 print("="*50)
13
14 soup = BeautifulSoup(sent_1500, 'lxml')
15 text = soup.get_text()
16 print(text)
17 print("="*50)
18
19 soup = BeautifulSoup(sent_4900, 'lxml')
20 text = soup.get_text()
21 print(text)
```

I can remember seeing the show when it aired on television years ago, when I was a child. My sister later bought me the LP (which I have to this day, I'm thirty something). I used this series of books & songs when I did my student teaching for preschoolers & turned the whole school on to it. I am now purchasing it on CD, along with the books for my children 5 & 2. The tradition lives on!

=====

Frank's Red Hot sauce is the best hot sauce there is. It's good on just about anything - even Chinese food! Seriously, try it on everything! It gives anything a really great flavor & kick.

=====

I purchased these on the special Friday sale where they were \$7.99 for 3 cans. At the store where I normally shop here one can is almost \$5 so this was a fantastic deal. I liked how they were lightly salted and not as salty as a lot of nuts are. I agree with the other review that they were mostly almonds. I'd say 60% almonds, 30% cashews and 10% macadamias but they were still quite good.

=====

This is the second time that I have ordered Belly Flops. They contain an excellent assortment of flavors. The odd shapes make them interesting, and the lack of the words "Jelly Belly" does not detract at all from the taste!

```
In [20]: 1 # https://stackoverflow.com/a/47091490/4084039
2 import re
3
4 def decontracted(phrase):
5     # specific
6     phrase = re.sub(r"won't", "will not", phrase)
7     phrase = re.sub(r"can't", "can not", phrase)
8
9     # general
10    phrase = re.sub(r"n't", " not", phrase)
11    phrase = re.sub(r"\ 're", " are", phrase)
12    phrase = re.sub(r"\ 's", " is", phrase)
13    phrase = re.sub(r"\ 'd", " would", phrase)
14    phrase = re.sub(r"\ 'll", " will", phrase)
15    phrase = re.sub(r"\ 't", " not", phrase)
16    phrase = re.sub(r"\ 've", " have", phrase)
17    phrase = re.sub(r"\ 'm", " am", phrase)
18    return phrase
```

```
In [21]: 1 sent_1500 = decontracted(sent_1500)
2 print(sent_1500)
3 print("="*50)
```

I purchased these on the special Friday sale where they were \$7.99 for 3 cans. At the store where I normally shop here one can is almost \$5 so this was a fantastic deal. I liked how they were lightly salted and not as salty as a lot of nuts are.

I agree with the other review that they were mostly almonds. I'd say 60% almonds, 30% cashews and 10% macadamias but they were still quite good.

=====

```
In [22]: 1 #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
2 sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
3 print(sent_0)
```

I can remember seeing the show when it aired on television years ago, when I was a child. My sister later bought me the LP (which I have to this day, I'm thirty something). I used this series of books & songs when I did my student teaching for preschoolers & turned the whole school on to it. I am now purchasing it on CD, along with the books for my children & The tradition lives on!

```
In [23]: 1 #remove spacial character: https://stackoverflow.com/a/5843547/4084039
2 sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
3 print(sent_1500)
```

I purchased these on the special Friday sale where they were 7 99 for 3 cans At the store where I normally shop here one can is almost 5 so this was a fantastic deal I liked how they were lightly salted and not as salty as alot of nuts are br br I agree with the other review that they were mostly almonds Id say 60 almonds 30 cashews and 10 macadamias but they were still quite good

```
In [24]: 1 # https://gist.github.com/sebleier/554280
2 # we are removing the words from the stop words List: 'no', 'nor', 'not'
3 # <br /><br /> ==> after the above steps, we are getting "br br"
4 # we are including them into stop words list
5 # instead of <br /> if we have <br/> these tags would have revmoved in the 1st step
6
7 stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "
8     "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
9     'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their'
10    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'tho
11    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', '
12    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', '
13    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before',
14    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again'
15    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'f
16    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', \
17    's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', '
18    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't",
19    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mus
20    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'were
21    'won', "won't", 'wouldn', "wouldn't"])
```



```
In [28]: 1 from sklearn.model_selection import train_test_split
2
3 X_train,X_test,Y_train,Y_test = train_test_split(preprocessed_reviews,final['Score'],test_size=0.3)
4 #splitting the dataset
5
6 print('Size of train dataset is:',len(X_train))#size of training dataset
7 print('Size of the test dataset is:',len(X_test))#size of test dataset
```

Size of train dataset is: 70000
Size of the test dataset is: 30000

```
In [29]: 1 from sklearn.model_selection import TimeSeriesSplit#importing for time series split
2 tscv = TimeSeriesSplit(n_splits=10)#time series split for the data
3 print(tscv)
```

TimeSeriesSplit(n_splits=10)

```
In [30]: 1 from sklearn.model_selection import TimeSeriesSplit
2 tscv = TimeSeriesSplit(n_splits=10)
3 for train, cv in tscv.split(X_train):
4     print('train data shape:',train.shape,'test data shape',cv.shape)
```

train data shape: (6370,) test data shape (6363,)
train data shape: (12733,) test data shape (6363,)
train data shape: (19096,) test data shape (6363,)
train data shape: (25459,) test data shape (6363,)
train data shape: (31822,) test data shape (6363,)
train data shape: (38185,) test data shape (6363,)
train data shape: (44548,) test data shape (6363,)
train data shape: (50911,) test data shape (6363,)
train data shape: (57274,) test data shape (6363,)
train data shape: (63637,) test data shape (6363,)

[5] Applying Random forest

[5.1] Applying Random Forests on BOW with RandomizedSearch CROSS VALIDATION, SET 1

```
In [51]: 1 bow_vect = CountVectorizer()#initiating the vectorizer
2
3 bow_vect.fit(X_train)#fitting data into vectorizer makes it learn all the vocablury
4
5 #transforming the data into training and test dataset
6 train_set_bow = bow_vect.transform(X_train)
7 test_set_bow = bow_vect.transform(X_test)
8
9 print('AFTER VECTORIZATION:')
10 print(train_set_bow.shape)
11 print(test_set_bow.shape)
12 print('Some of the feature names are:',bow_vect.get_feature_names()[ :10:-1])
13
```

AFTER VECTORIZATION:

(70000, 50810)

(30000, 50810)

```
In [52]: 1 from sklearn.ensemble import RandomForestClassifier
2 from sklearn.model_selection import RandomizedSearchCV
3
4 estimators = [10,50,100,250,450]#list of estimators that will be tuned
5 depths = [3,9,11,15,50]#tuning depth to avoid overfitting and underfitting
6
7 params = {'max_depth':depths,'n_estimators':estimators}#for passing as argument
8
9 tscv = TimeSeriesSplit(n_splits = 5)#initiating 5 time series splits for cross validation
10
11 model = RandomizedSearchCV(RandomForestClassifier(bootstrap = True,criterion = 'gini',max_features = 'auto')
12                             return_train_score=True,n_jobs = -1)
13 #fitting the classifier
14 #fitting the parameter distribution
15
16
17 model.fit(train_set_bow,Y_train)
18
19
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 10.4min
[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 30.5min finished
```

```
Out[52]: RandomizedSearchCV(cv=TimeSeriesSplit(n_splits=5), error_score='raise',
    estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=None, max_features='auto', max_leaf_nodes=None,
    min_impurity_split=1e-07, min_samples_leaf=1,
    min_samples_split=2, min_weight_fraction_leaf=0.0,
    n_estimators=10, n_jobs=1, oob_score=False, random_state=None,
    verbose=0, warm_start=False),
    fit_params={}, iid=True, n_iter=10, n_jobs=-1,
    param_distributions={'max_depth': [3, 9, 11, 15, 50], 'n_estimators': [10, 50, 100, 250, 450]},
    pre_dispatch='2*n_jobs', random_state=None, refit=True,
    return_train_score=True, scoring='roc_auc', verbose=1)
```

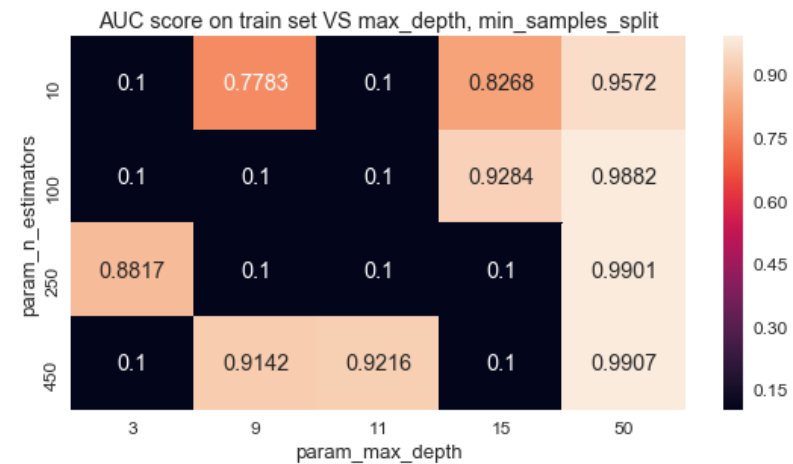

```
In [53]: 1 savetofile(model,'model_bow_rf')#saving the model
```

```
In [54]: 1 model_bow_rf = openfromfile('model_bow_rf')#retreiving the model
```

```
In [55]: 1 # as we have two hyperparameters to tune so we will plot heatmap and to show hyperparameters giving maximum
2 def plots(model):#function for plotting heatmaps
3
4     print('Best Hyperparameters are:',model.best_params_)
5     df = pd.DataFrame(model.cv_results_)#saving into the dataframe
6     results = df.groupby(['param_n_estimators','param_max_depth']).min().unstack()[['mean_test_score',
7                                                                                       'mean_train_score']]
8     #groupby by number of estimators and maximum depth and unstacking mean train and test score
9
10    results = results.fillna(0.1)#imputing all null values by 0.1
11
12    sns.set(font_scale = 1.2)
13    fig, ax = plt.subplots(figsize=(20,10))#setting the font size
14    plt.subplot(2,2,1)
15    title_test = 'AUC score on test set VS max_depth, min_samples_split'
16    fmt = 'png'
17    sns.heatmap(results.mean_test_score, annot=True, fmt='.4g');#heatmap for test score
18    plt.title(title_test);
19    #plt.savefig('{title_test}.{fmt}', format=fmt, dpi=300);
20
21    plt.subplot(2,2,2)
22    title_train = 'AUC score on train set VS max_depth, min_samples_split'
23    fmt = 'png'
24    sns.heatmap(results.mean_train_score, annot=True, fmt='.4g');#heatmap for train score
25    plt.title(title_train);
26    #plt.savefig('{title_train}.{fmt}', format=fmt, dpi=300);
```

In [56]: 1 plots(model_bow_rf)

Best Hyperparameters are: {'n_estimators': 450, 'max_depth': 50}



In [64]:

```

1  from sklearn.ensemble import RandomForestClassifier
2  from sklearn.metrics import roc_auc_score
3  from sklearn.metrics import roc_curve
4  from sklearn.metrics import confusion_matrix
5  from sklearn.metrics import precision_score
6  from sklearn.metrics import recall_score
7  from sklearn.metrics import f1_score
8  from wordcloud import WordCloud
9  #from xgboost import XGBClassifier
10
11
12  #1.Function for calculating the test and train Area under curve after fitting with right hyperparameters
13  def auc(depth,estimator,train_set,test_set):
14      tree_optimal = RandomForestClassifier(bootstrap = True,criterion = 'gini',max_depth = depth,n_estimators
15      tree_optimal.fit(train_set,Y_train)
16      pred_tr = tree_optimal.predict(train_set)# predicting all the classes for test dataset for confusion mat
17      pred_test = tree_optimal.predict(test_set)#predicting all the classes for train dataset for confusin mat
18
19      train_pred_proba = tree_optimal.predict_proba(train_set)[:,-1]
20      test_pred_proba = tree_optimal.predict_proba(test_set)[:,-1]
21      #predict_proba gives the probability of a particular data point belonging to the specified class
22
23      train_auc = roc_auc_score(Y_train,train_pred_proba)
24      test_auc = roc_auc_score(Y_test,test_pred_proba)
25      print('AUC on train data is:',train_auc)
26      print('AUC on test data is:',test_auc)
27      print("*****\n")
28      return train_auc,test_auc,train_pred_proba,test_pred_proba,pred_tr,pred_test
29
30
31  #*****
32
33  #2.Function for plotting the roc curve
34  def curve(train_pred,test_pred ):
35      fpr_tr, tpr_tr, _ = roc_curve(Y_train,train_pred)
36      fpr_test, tpr_test, _ = roc_curve(Y_test,test_pred)
37      #calculating the fpr,tpr and thresholds for each training and test dataset
38      auc_train = roc_auc_score(Y_train,train_pred)
39      auc_test = roc_auc_score(Y_test, test_pred)
40      sns.set_style('darkgrid')
41      plt.figure(figsize=(8,8))
42      plt.plot(np.linspace(0,1,100),np.linspace(0,1,100),"g--")#this plots the roc curve for AUC = 0.5

```

```

43 plt.plot(fpr_tr, tpr_tr, 'r', linewidth=2, label="train auc="+str(auc_train))
44 plt.plot(fpr_test, tpr_test, 'b', linewidth=1, label=" test auc="+str(auc_test))
45 plt.xlabel('False positive rate(1-specificity)', fontsize=18)
46 plt.ylabel('True positive rate(sensitivity)', fontsize=18)
47 plt.title('Reciever operating characteristics curve', fontsize=18)
48 plt.legend(loc='best')
49 plt.show()
50 print('*****\n')
51
52 #*****
53
54 #3.Function for calculating F1, precision and recall
55 def metrics(pred):
56     print('scores on test data are:\n')
57     # calculating the precison score
58     print('precision score is {}'.format(precision_score(Y_test, pred)))
59     #calculating the recall score
60     print('\nrecall_score is {}'.format(recall_score(Y_test, pred)))
61     #calculating the f1 score
62     print('\nf1 score is {}\n'.format(f1_score(Y_test, pred)))
63     print('*****\n")
64
65 #*****
66
67 #4.Function for plotting the confusion matrix
68 def plot_confusion_matrix(test_y, predict_y):
69     C = confusion_matrix(test_y, predict_y)
70
71     A = (((C.T)/(C.sum(axis=1))).T)#for recall matrix
72
73     B =(C/C.sum(axis=0))#for precision matrix
74     plt.figure(figsize=(20,4))
75
76     labels = [0,1]
77     # representing A in heatmap format
78     cmap=sns.light_palette("blue")
79     plt.subplot(1, 3, 1)
80     sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
81     plt.xlabel('Predicted Class')
82     plt.ylabel('Original Class')
83     plt.title("Confusion matrix")
84
85     plt.subplot(1, 3, 2)

```



```

86 sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
87 plt.xlabel('Predicted Class')
88 plt.ylabel('Original Class')
89 plt.title("Precision matrix")
90
91 plt.subplot(1, 3, 3)
92 # representing B in heatmap format
93 sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
94 plt.xlabel('Predicted Class')
95 plt.ylabel('Original Class')
96 plt.title("Recall matrix")
97
98 plt.show()
99 #*****
100 def features_and_wc(depth,estimator,train_set,vectorizer):
101     text = "" #saving strings
102     tree_optimal = RandomForestClassifier(criterion = 'gini',max_depth = depth,n_estimators = estimator)
103     #tree_optimal = XGBClassifier(learning_rate = 0.1,booster = "gbtree",max_depth = depth,n_estimators = es
104     tree_optimal.fit(train_set,Y_train)
105     features = tree_optimal.feature_importances_
106     indices = np.argsort(features)[::-1]
107     feature_names = vectorizer.get_feature_names()
108     print('TOP 20 important features which gives maximum information gain on splitting are:\n')
109     for i in (indices[0:20]):
110         text = text + " " + feature_names[i]
111         print("%s\t -->\t%f" %(feature_names[i],features[i]))
112
113     wordcloud = WordCloud(width=1500, height=600,stopwords = stopwords).generate(text)
114     # plot the WordCloud image
115     plt.figure(figsize = (30,8))
116     plt.imshow(wordcloud, interpolation="bilinear")
117     plt.axis("off")
118     plt.margins(x=0, y=0)
119     plt.show()
120
121

```

Here 0.1 signifies that set of parameters from the list were not used for tuning in randomized search cv

```
In [70]: 1 from IPython.core.display import display, HTML
          2 display(HTML("<style>.container { width:100% !important; }</style>"))
```

ROC and Confusion Matrix

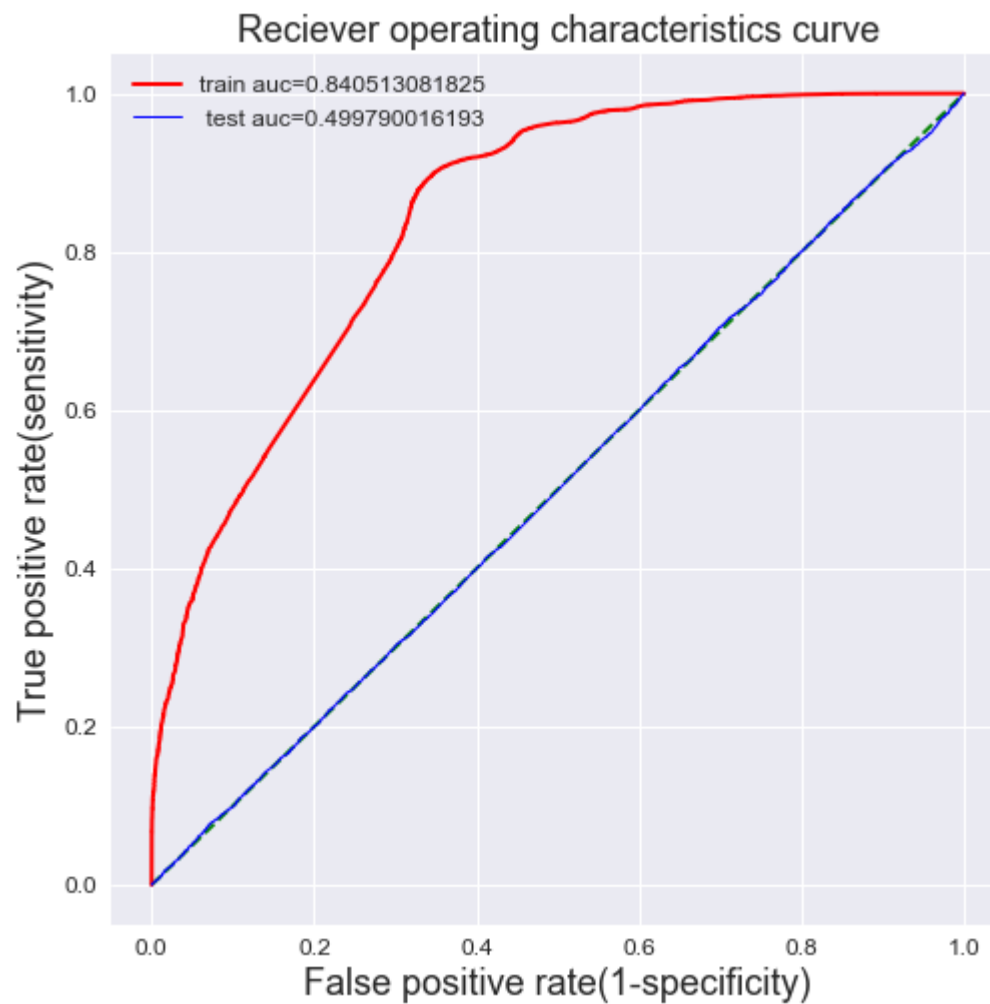
```
In [65]: 1 best_depth_bow_rf = model_bow_rf.best_params_['max_depth']
          2 best_n_bow_rf = model_bow_rf.best_params_['n_estimators']
```

In [66]:

```
1  '''AUC ON TEST DATA'''
2  train_auc_BOW_rf, test_auc_BOW_rf, train_pred_proba, test_pred_proba, train_pred, test_pred = auc(best_depth_bow_
3
4  '''PLOTING THE ROC CURVE'''
5  curve(train_pred_proba, test_pred_proba)
6
7  '''Precision, recall and f1 score'''
8  metrics(test_pred)
9
10
11
12  '''Plotting the confusion matrix'''
13  print('Confusin matrix on train data')
14  plot_confusion_matrix(Y_train, train_pred)
15  print('Confusion matrix for test data')
16  plot_confusion_matrix(Y_test, test_pred)
```

AUC on train data is: 0.840513081825

AUC on test data is: 0.499790016193




scores on test data are:

precison score is 0.8433


recall_score is 1.0

f1 score is 0.9149894211468561

Confusin matrix on train data



Confusion matrix for test data




Feature Importance and WORDCLOUD

In [69]:

```
1 """Most important features"""  
2 features_and_wc(best_depth_bow_rf,best_n_bow_rf,train_bow,bow_vect)
```

TOP 20 important features which gives maximum information gain on splitting are:

```
nota      --> 0.002619  
cuantro   --> 0.001636  
baily     --> 0.001557  
resulyed  --> 0.001541  
ligtly    --> 0.001520  
cassava   --> 0.001496  
graphic   --> 0.001443  
goth      --> 0.001434  
procyanidins --> 0.001418  
grimaced  --> 0.001315  
tap       --> 0.001277  
dependents --> 0.001249  
nanak     --> 0.001243  
omission  --> 0.001219  
smothering --> 0.001210  
girft     --> 0.001183  
flo       --> 0.001142  
chinsy    --> 0.001134  
woould    --> 0.001129  
nodding   --> 0.001110
```



[5.2] Applying Random Forests on TFIDF, SET 2

```
In [0]: 1 # Please write all the code with proper documentation
2 tfidf_vect = TfidfVectorizer(ngram_range = (1,2),min_df = 10)
3 #min_df signifies minimum number of times a word must occur in corpus for consideration
4 #ngram_range tells about the unigram and bigram
5 tfidf_vect.fit(X_train)
6 train_set = tfidf_vect.transform(X_train)
7 test_set = tfidf_vect.transform(X_test)
8
9 print('after vectorization training set:',train_set.shape)
10 print('after vectorization test set:',test_set.shape)
11
```

after vectorization training set: (70000, 40504)

after vectorization test set: (30000, 40504)

```
In [0]: 1 savetofile(train_set,'train_tfidf')#saving to file for future use
        2 savetofile(test_set,'test_tfidf')
```

```
In [71]: 1 train_tfidf = openfromfile('computed/train_tfidf')
        2 test_tfidf = openfromfile('computed/test_tfidf')
```

```
In [0]: 1
        2 estimators = [10,50,100,250,450]#list of estimators that will be tuned
        3 depths = [3,9,11,15,50]#tuning depth to avoid overfitting and underfitting
        4
        5 params = {'max_depth':depths,'n_estimators':estimators}#for passing as argument
        6
        7 tscv = TimeSeriesSplit(n_splits = 5)#initiating 5 time series splits for cross validation
        8
        9 model = RandomizedSearchCV(RandomForestClassifier(bootstrap = True,criterion = 'gini',max_features = 'auto')
       10                           return_train_score=True,n_jobs = -1)
       11 #fitting the classifier
       12 #fitting the parameter distribution
       13
       14
       15 model.fit(train_tfidf,Y_train)
       16
       17
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 4.5min
[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 5.0min finished
```

```
Out[50]: RandomizedSearchCV(cv=TimeSeriesSplit(n_splits=5), error_score='raise',
                             estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             min_impurity_split=1e-07, min_samples_leaf=1,
                             min_samples_split=2, min_weight_fraction_leaf=0.0,
                             n_estimators=10, n_jobs=1, oob_score=False, random_state=None,
                             verbose=0, warm_start=False),
                             fit_params={}, iid=True, n_iter=10, n_jobs=-1,
                             param_distributions={'max_depth': [3, 9, 11, 15, 50], 'n_estimators': [10, 50, 100, 250, 450]},
                             pre_dispatch='2*n_jobs', random_state=None, refit=True,
                             return_train_score=True, scoring='roc_auc', verbose=1)
```

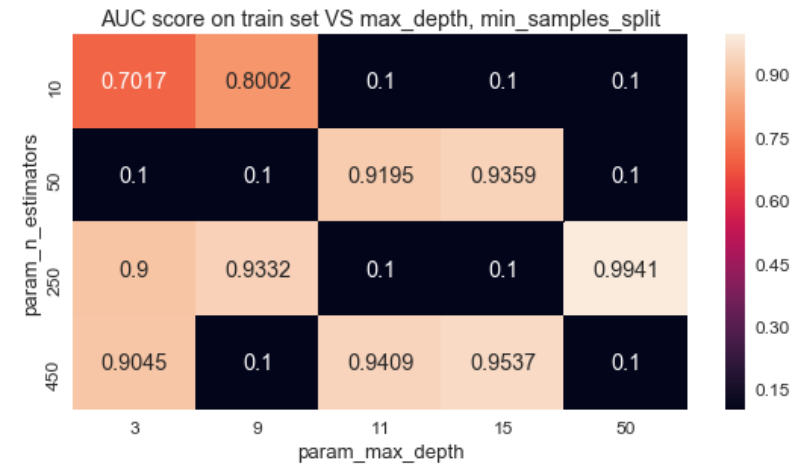



```
In [0]: 1 savetofile(model, 'model_tfidf_rf')
```

```
In [72]: 1 model_tfidf_rf = openfromfile('model_tfidf_rf')
```

```
In [73]: 1 plots(model_tfidf_rf)
```

Best Hyperparameters are: {'n_estimators': 250, 'max_depth': 50}



```
In [74]: 1 best_depth_tfidf_rf = model_tfidf_rf.best_params_['max_depth'] #best depth of the decision trees
          2 best_n_tfidf_rf = model_tfidf_rf.best_params_['n_estimators'] #best number of decision tress
```


ROC and Confusion Matrix

In [75]:

```
1  '''AUC ON TEST DATA'''
2  train_auc_tfidf_rf,test_auc_tfidf_rf,train_pred_proba,test_pred_proba,train_pred,test_pred = auc(best_depth_
3
4  '''PLOTING THE ROC CURVE'''
5  curve(train_pred_proba,test_pred_proba)
6
7  '''Precision,recall and f1 score'''
8  metrics(test_pred)
9
10
11  '''Plotting the confusion matrix'''
12  print('Confusin matrix on train data')
13  plot_confusion_matrix(Y_train,train_pred)
14  print('Confusion matrix for test data')
15  plot_confusion_matrix(Y_test,test_pred)
```

AUC on train data is: 0.84572698547

AUC on test data is: 0.495126376182




scores on test data are:

precison score is 0.8433


recall_score is 1.0

f1 score is 0.9149894211468561

Confusin matrix on train data



Confusion matrix for test data



Feature importance and WORDCLOUD

```
In [0]: 1 ""Most important features""  
2 features_and_wc(best_depth_tfidf_rf,best_n_tfidf_rf,train_tfidf,tfidf_vect)#for generating the wordcloud
```

TOP 20 important features which gives maximum information gain on splitting are:

not	-->	0.016877
great	-->	0.011052
disappointed	-->	0.010158
not buy	-->	0.009911
awful	-->	0.009760
worst	-->	0.009393
horrible	-->	0.008511
money	-->	0.008419
waste money	-->	0.008219
bad	-->	0.008164
waste	-->	0.007668
not recommend	-->	0.006845
disappointing	-->	0.006349
terrible	-->	0.006257
would not	-->	0.006193
best	-->	0.006062
threw	-->	0.005787
return	-->	0.005548
love	-->	0.005523
refund	-->	0.005472



[5.3] Applying Random Forests on AVG W2V and TFIDF W2V, SET 3

```
In [0]: 1 s_train = []
        2 for sent in X_train:
        3     s_train.append(sent.split())
        4 #preparing the training data for word to vector vectorization
        5
        6 s_test = []
        7 for sent in X_test:
        8     s_test.append(sent.split())
        9 #preparing the test data for word to vector fatorization
```

```
In [0]: 1 # this line of code trains the w2v model on the give list of sentences
2
3 w2v_model=Word2Vec(s_train,min_count=5,size=50, workers=4)# min_count = 5 considers only words that occurred
4
5 w2v_words = list(w2v_model.wv.vocab)
6 print("number of words that occurred minimum 5 times ",len(w2v_words))
7
```

number of words that occurred minimum 5 times 16003

```
In [0]: 1 print('sample words :',w2v_words[0:100])
```

sample words : ['four', 'cats', 'three', 'absolutely', 'love', 'treats', 'one', 'sometimes', 'accepts', 'however', 'ate', 'treat', 'would', 'not', 'touch', 'sniffed', 'walked', 'away', 'lovers', 'based', 'reaction', 'say', 'chance', 'cat', 'liking', 'pleased', 'delivery', 'product', 'expected', 'branch', 'mints', 'purchased', 'years', 'best', 'need', 'change', 'packaging', 'twist', 'ends', 'pressed', 'ease', 'taking', 'package', 'problem', 'made', 'mexico', 'usa', 'kraft', 'cheese', 'fill', 'blank', 'products', 'predictable', 'well', 'suited', 'young', 'palates', 'followed', 'box', 'recipe', 'two', 'changes', 'instead', 'chicken', 'flaked', 'oz', 'leftover', 'salmon', 'also', 'melted', 'half', 'dozen', 'slices', 'american', 'stronger', 'flavor', 'kid', 'friendly', 'unlike', 'frozen', 'broccoli', 'prominent', 'florets', 'tiny', 'dehydrated', 'reconstitute', 'much', 'failed', 'contribute', 'depending', 'eaters', 'may', 'consider', 'pro', 'con', 'hard', 'see', 'anybody', 'better', 'boxed']

[5.3.1] Average word to vector

```

1  #computing average word to vector for training data
2  train_set = [] # the avg-w2v for each sentence/review is stored in this list
3  for sent in tqdm(s_train):
4      sent_vec = np.zeros(50)
5      cnt_words = 0; # num of words with a valid vector in the sentence/review
6      for word in sent: #
7          if word in w2v_words:
8              vec = w2v_model.wv[word]
9              sent_vec += vec
10             cnt_words += 1
11     if cnt_words != 0:
12         sent_vec /= cnt_words
13     train_set.append(sent_vec)
14
15     print(len(train_set))#number of data points

```

70000


```

In [0]: 1 estimators = [10,50,100,250,450]#list of estimators that will be tuned
2 depths = [3,9,11,15,50]#tuning depth to avoid overfitting and underfitting
3 params = {'max_depth':depths,'n_estimators':estimators}#for passing as argument
4 tscv = TimeSeriesSplit(n_splits = 5)#initiating 5 time series splits for cross validation
5 model = RandomizedSearchCV(RandomForestClassifier(bootstrap = True,criterion = 'gini',max_features = 'auto')
6                               return_train_score=True,n_jobs = -1)
7 #fitting the classifier
8 #fitting the parameter distribution
9 model.fit(train_avgw2v,Y_train)
10
11

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```

[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 7.4min
[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 11.2min finished

```

```

Out[65]: RandomizedSearchCV(cv=TimeSeriesSplit(n_splits=5), error_score='raise',
    estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=None, max_features='auto', max_leaf_nodes=None,
    min_impurity_split=1e-07, min_samples_leaf=1,
    min_samples_split=2, min_weight_fraction_leaf=0.0,
    n_estimators=10, n_jobs=1, oob_score=False, random_state=None,
    verbose=0, warm_start=False),
    fit_params={}, iid=True, n_iter=10, n_jobs=-1,
    param_distributions={'max_depth': [3, 9, 11, 15, 50], 'n_estimators': [10, 50, 100, 250, 450]},
    pre_dispatch='2*n_jobs', random_state=None, refit=True,
    return_train_score=True, scoring='roc_auc', verbose=1)

```

```

In [0]: 1 savetofile(model,'model_avgw2v_rf')

```

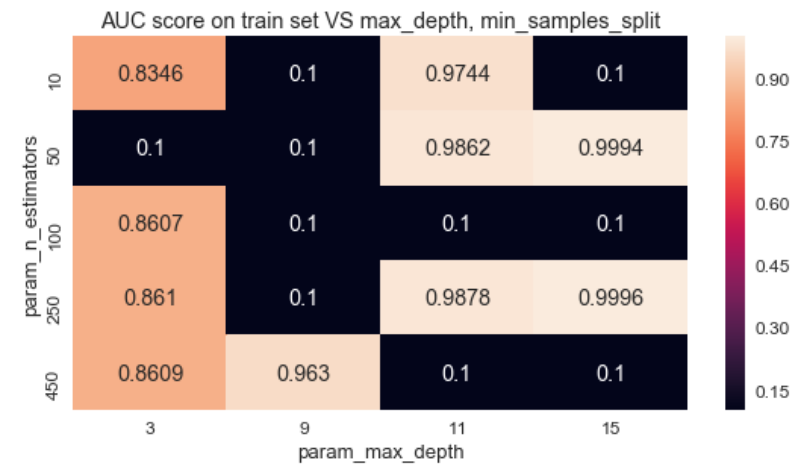

```

In [0]: 1 model_avgw2v_rf = openfromfile('model_avgw2v_rf')

```

```
In [0]: 1 plots(model_avgw2v_rf)
```

Best Hyperparameters are: {'n_estimators': 250, 'max_depth': 15}




ROC and Confusion Matrix

```
In [0]: 1 '''AUC ON TEST DATA'''
2 train_auc_avgw2v_rf,test_auc_avgw2v_rf,train_pred_proba,test_pred_proba,train_pred,test_pred = auc(best_dept
3
4 '''PLOTING THE ROC CURVE'''
5 curve(train_pred_proba,test_pred_proba)
6
7 '''Precision,recall and f1 score'''
8 metrics(test_pred)
9
10
11 '''Plotting the confusion matrix'''
12 print('Confusin matrix on train data')
13 plot_confusion_matrix(Y_train,train_pred)
14 print('Confusion matrix for test data')
15 plot_confusion_matrix(Y_test,test_pred)
```

AUC on train data is: 0.998923590696

AUC on test data is: 0.903633184442




scores on test data are:

precison score is 0.8898365569909357


recall_score is 0.9848724227822103

f1 score is 0.9349456317960255

Confusin matrix on train data



Confusion matrix for test data



[5.3.2] Applying Random Forests on TFIDF W2V, SET 4

```
In [0]: 1 vect_tfidfw2v = TfidfVectorizer()#initializing the tfidf vectorizer
        2
        3 tf_idf = vect_tfidfw2v.fit_transform(X_train)#fitting the training data
        4 dictionary = dict(zip(vect_tfidfw2v.get_feature_names(), list(vect_tfidfw2v.idf_)))#zipping both of the feat
```

```
In [0]: 1 import itertools
        2 dict(itertools.islice(dictionary.items(),20))
        3 #printing first 20 elements of the dictionary
```

```
Out[73]: {'aa': 9.5172074770284834,
          'aaa': 10.546826894209641,
          'aaaa': 11.463117626083797,
          'aaaaa': 11.463117626083797,
          'aaaaaa': 11.463117626083797,
          'aaaaaaaa': 11.463117626083797,
          'aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa': 11.463117626083797,
          'aaaaahhhhyaaaaaa': 11.463117626083797,
          'aaaaaand': 11.463117626083797,
          'aaaaah': 11.463117626083797,
          'aaaah': 11.057652517975633,
          'aaaand': 11.463117626083797,
          'aaaannnnddd': 11.463117626083797,
          'aaagh': 11.463117626083797,
          'aaah': 10.769970445523851,
          'aaahhh': 11.463117626083797,
          'aaahs': 11.463117626083797,
          'aadp': 11.463117626083797,
          'aafco': 10.076823264963906,
          'aafes': 10.769970445523851}
```

```
In [0]: 1 tfidf_feat = vect_tfidfw2v.get_feature_names() # tfidf words/col-names
        2 print(tfidf_feat[:100])
```

```
['aa', 'aaa', 'aaaa', 'aaaaa', 'aaaaaa', 'aaaaaaaa', 'aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa', 'aaaaa
ahhhhyaaaaaa', 'aaaaaand', 'aaaaah', 'aaaah', 'aaaand', 'aaaannnnddd', 'aaagh', 'aaah', 'aaahhh', 'aaahs', 'aa
dp', 'aafco', 'aafes', 'aafter', 'aah', 'aahhhs', 'aahing', 'aalouisiana', 'aanother', 'aap', 'aarrgh', 'aart
i', 'ab', 'aback', 'abandon', 'abandoned', 'abandoning', 'abbey', 'abbott', 'abby', 'abc', 'abd', 'abdomen', 'a
bdominal', 'abduct', 'aber', 'abhor', 'abhorrent', 'abide', 'abiding', 'abietate', 'abig', 'abigirl', 'abiiit
y', 'abilities', 'ability', 'abit', 'abj', 'abject', 'abjectly', 'ablation', 'ablaze', 'able', 'abletate', 'abl
eto', 'ablsolutely', 'ably', 'abmout', 'abnormal', 'abnormalities', 'abnormally', 'abnoxious', 'abominable', 'a
bominably', 'abomination', 'abondant', 'aboout', 'aboslutely', 'abosolutely', 'abot', 'abottle', 'abou', 'aboun
d', 'aboutamazon', 'aboutthis', 'abouti', 'aboutits', 'aboutletting', 'aboutthis', 'aboutwhat', 'abouy', 'aboveo
ur', 'abover', 'abp', 'abra', 'abraham', 'abrasion', 'abrasions', 'abreakfast', 'abreast', 'abbreviating', 'abri
dged', 'abroad']
```



```

In [0]: 1 estimators = [10,50,100,250,450]#list of estimators that will be tuned
        2 depths = [3,9,11,15,50]#tuning depth to avoid overfitting and underfitting
        3 params = {'max_depth':depths,'n_estimators':estimators}#for passing as argument
        4 tscv = TimeSeriesSplit(n_splits = 5)#initiating 5 time series splits for cross validation
        5 model = RandomizedSearchCV(RandomForestClassifier(bootstrap = True,criterion = 'gini',max_features = 'auto')
        6                               return_train_score=True,n_jobs = -1)
        7 #fitting the classifier
        8 #fitting the parameter distribution
        9 model.fit(train_tfidf2v,Y_train)
       10
       11

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```

[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 11.0min
[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 14.7min finished

```

```

Out[81]: RandomizedSearchCV(cv=TimeSeriesSplit(n_splits=5), error_score='raise',
    estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=None, max_features='auto', max_leaf_nodes=None,
    min_impurity_split=1e-07, min_samples_leaf=1,
    min_samples_split=2, min_weight_fraction_leaf=0.0,
    n_estimators=10, n_jobs=1, oob_score=False, random_state=None,
    verbose=0, warm_start=False),
    fit_params={}, iid=True, n_iter=10, n_jobs=-1,
    param_distributions={'max_depth': [3, 9, 11, 15, 50], 'n_estimators': [10, 50, 100, 250, 450]},
    pre_dispatch='2*n_jobs', random_state=None, refit=True,
    return_train_score=True, scoring='roc_auc', verbose=1)

```

```

In [0]: 1 savetofile(model,'model_tfidf2v_rf')

```



```

In [0]: 1 model_tfidf2v_rf = openfromfile('model_tfidf2v_rf')

```

```
In [0]: 1 plots(model_tfidf2v_rf)
```

Best Hyperparameters are: {'n_estimators': 450, 'max_depth': 50}




```
In [0]: 1 best_depth_tfidf2v_rf = model_tfidf2v_rf.best_params_['max_depth']
        2 best_n_tfidf2v_rf = model_tfidf2v_rf.best_params_['n_estimators']
```

```
In [0]: 1 '''AUC ON TEST DATA'''
2 train_auc_tfidf2v_rf, test_auc_tfidf2v_rf, train_pred_proba, test_pred_proba, train_pred, test_pred = auc(best_
3
4 '''PLOT THE ROC CURVE'''
5 curve(train_pred_proba, test_pred_proba)
6
7 '''Precision, recall and f1 score'''
8 -metrics(test_pred)
9
10
11 '''Plotting the confusion matrix'''
12 print('Confusion matrix on train data')
13 plot_confusion_matrix(Y_train, train_pred)
14 print('Confusion matrix for test data')
15 plot_confusion_matrix(Y_test, test_pred)
```

AUC on train data is: 0.999805975597

AUC on test data is: 0.883013839274




scores on test data are:

precison score is 0.8792848595762652


recall_score is 0.9868078047239118

f1 score is 0.9299486339611404

Confusin matrix on train data



Confusion matrix for test data



[6] Applying GBDT using XGBOOST

[6.1] Applying XGBOOST on BOW, **SET 1**

```

In [0]: 1 from xgboost import XGBClassifier
        2 from sklearn.model_selection import RandomizedSearchCV
        3 from sklearn.model_selection import TimeSeriesSplit
        4
        5
        6
        7 estimators = [10,50,100,200,450]#List of estimators that will be tuned
        8 depths = [3,9,11,13,15]#tuning depth to avoid overfitting and underfitting
        9 params = {'max_depth':depths,'n_estimators':estimators}#for passing as argument
       10
       11 tscv = TimeSeriesSplit(n_splits = 5)
       12 model = RandomizedSearchCV(XGBClassifier(booster = 'gbtree',learning_rate = 0.1),param_distributions = param
       13 model.fit(train_bow,Y_train)#fitting the model
       14
       15

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n_jobs=1)]: Done 50 out of 50 | elapsed: 73.7min finished

```

Out[60]: RandomizedSearchCV(cv=TimeSeriesSplit(max_train_size=None, n_splits=5),
                             error_score='raise-deprecating',
                             estimator=XGBClassifier(base_score=0.5, booster='gbtree',
                                                         colsample_bylevel=1,
                                                         colsample_bynode=1,
                                                         colsample_bytree=1, gamma=0,
                                                         learning_rate=0.1, max_delta_step=0,
                                                         max_depth=3, min_child_weight=1,
                                                         missing=None, n_estimators=100,
                                                         n_jobs=1, nthread=None,
                                                         objective='binary:logistic',
                                                         random_state=0, reg_alpha=0,
                                                         reg_lambda=1, scale_pos_weight=1,
                                                         seed=None, silent=None, subsample=1,
                                                         verbosity=1),
                             iid='warn', n_iter=10, n_jobs=None,
                             param_distributions={'max_depth': [3, 9, 11, 13, 15],
                                                    'n_estimators': [10, 50, 100, 200,
                                                                    450]}},
                             pre_dispatch='2*n_jobs', random_state=None, refit=True,
                             return_train_score=True, scoring=None, verbose=1)



```

```
In [0]: 1 savetofile(model, 'model_bow_gb')
```

```
In [0]: 1 model_bow_gb = openfromfile('model_bow_gb')
```

```
In [0]: 1 plots(model_bow_gb)
```

Best Hyperparameters are: {'n_estimators': 100, 'max_depth': 3}




```
In [0]: 1 best_depth_bow_gb = model_bow_gb.best_params_['max_depth']
        2 best_n_bow_gb = model_bow_gb.best_params_['n_estimators']
```



```
In [0]: 1 '''AUC ON TEST DATA'''
2 train_auc_bow_gb,test_auc_bow_gb,train_pred_proba,test_pred_proba,train_pred,test_pred = auc(best_depth_bow_
3
4 '''PLOTING THE ROC CURVE'''
5 curve(train_pred_proba,test_pred_proba)
6
7 '''Precision,recall and f1 score'''
8 metrics(test_pred)
9
10
11 '''Plotting the confusion matrix'''
12 print('Confusin matrix on train data')
13 plot_confusion_matrix(Y_train,train_pred)
14 print('Confusion matrix for test data')
15 plot_confusion_matrix(Y_test,test_pred)
```

AUC on train data is: 0.589092490071771

AUC on test data is: 0.5008964517846474




scores on test data are:

precison score is 0.8429


recall_score is 1.0

f1 score is 0.9147539204514623

Confusin matrix on train data




Confusion matrix for test data



```
In [0]: 1 """Most important features"""  
2 features_and_wc(best_depth_bow_gb,best_n_bow_gb,train_bow,bow_vect)#for generating the wordcloud
```

TOP 20 important features which gives maximum information gain on splitting are:

```
dorms      --> 0.007687  
fragments  --> 0.007296  
packedand  --> 0.007185  
flavorpineappleblue --> 0.007135  
steveia    --> 0.006729  
wake       --> 0.006622  
terra      --> 0.006486  
wiskas     --> 0.006415  
exited     --> 0.006380  
example    --> 0.006345  
stover     --> 0.005949  
souther    --> 0.005792  
returne    --> 0.005623  
pricefor   --> 0.005567  
slackened  --> 0.005395  
tahari     --> 0.005368  
disipate   --> 0.005320  
verify     --> 0.005293  
ribeyes    --> 0.005286  
unheard    --> 0.005201
```



[6.2] Applying XGBOOST on TFIDF, SET 2

```

In [0]: 1
2 estimators = [10,50,100,200,450]#list of estimators that will be tuned
3 depths = [3,9,11,13,15]#tuning depth to avoid overfitting and underfitting
4 params = {'max_depth':depths,'n_estimators':estimators}#for passing as argument
5
6 tscv = TimeSeriesSplit(n_splits = 5)
7 model = RandomizedSearchCV(XGBClassifier(booster = 'gbtree',learning_rate = 0.1),param_distributions = param
8 model.fit(train_tfidf,Y_train)#fitting the model
9
10

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n_jobs=1)]: Done 50 out of 50 | elapsed: 96.0min finished

```

Out[76]: RandomizedSearchCV(cv=TimeSeriesSplit(max_train_size=None, n_splits=5),
    error_score='raise-deprecating',
    estimator=XGBClassifier(base_score=0.5, booster='gbtree',
        colsample_bylevel=1,
        colsample_bynode=1,
        colsample_bytree=1, gamma=0,
        learning_rate=0.1, max_delta_step=0,
        max_depth=3, min_child_weight=1,
        missing=None, n_estimators=100,
        n_jobs=1, nthread=None,
        objective='binary:logistic',
        random_state=0, reg_alpha=0,
        reg_lambda=1, scale_pos_weight=1,
        seed=None, silent=None, subsample=1,
        verbosity=1),
    iid='warn', n_iter=10, n_jobs=None,
    param_distributions={'max_depth': [3, 9, 11, 13, 15],
        'n_estimators': [10, 50, 100, 200,
            450]},
    pre_dispatch='2*n_jobs', random_state=None, refit=True,
    return_train_score=True, scoring=None, verbose=1)

```

```



In [0]: 1 savetofile(model,'model_tfidf_gb')

```

```
In [0]: 1 model_tfidf_gb = openfromfile('model_tfidf_gb')
```

```
In [0]: 1 plots(model_tfidf_gb)#plotting for the model
```

Best Hyperparameters are: {'n_estimators': 200, 'max_depth': 3}



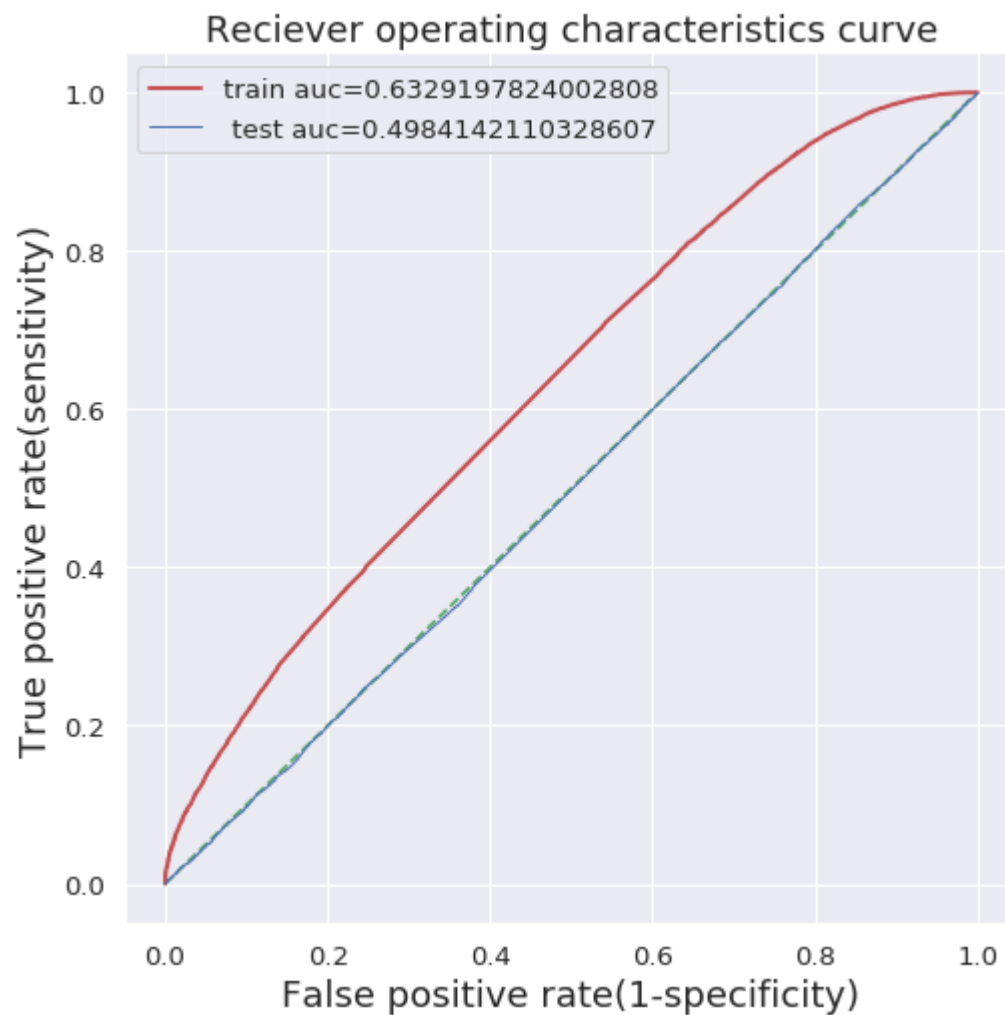
```
In [0]: 1
```

```
In [0]: 1 best_depth_tfidf_gb = model_tfidf_gb.best_params_['max_depth']
2 best_n_tfidf_gb = model_tfidf_gb.best_params_['n_estimators']
```

```
In [0]: 1 '''AUC ON TEST DATA'''
2 train_auc_tfidf_gb,test_auc_tfidf_gb,train_pred_proba,test_pred_proba,train_pred,test_pred = auc(best_depth_
3
4 '''PLOTING THE ROC CURVE'''
5 curve(train_pred_proba,test_pred_proba)
6
7 '''Precision,recall and f1 score'''
8 metrics(test_pred)
9
10
11 '''Plotting the confusion matrix'''
12 print('Confusin matrix on train data')
13 plot_confusion_matrix(Y_train,train_pred)
14 print('Confusion matrix for test data')
15 plot_confusion_matrix(Y_test,test_pred)
```

AUC on train data is: 0.6329197824002808

AUC on test data is: 0.4984142110328607




scores on test data are:

precison score is 0.8386666666666667


recall_score is 1.0

f1 score is 0.9122552574329225

Confusin matrix on train data



Confusion matrix for test data



[6.3] Applying XGBOOST on AVG W2V, SET 3

```
In [42]: 1 from sklearn.model_selection import RandomizedSearchCV
2 from xgboost import XGBClassifier
3 estimators = [10,50,100,200,450]#list of estimators that will be tuned
4 depths = [3,9,11,13,15]#tuning depth to avoid overfitting and underfitting
5 params = {'max_depth':depths,'n_estimators':estimators}#for passing as argument
6
7 tscv = TimeSeriesSplit(n_splits = 5)
8 model = RandomizedSearchCV(XGBClassifier(booster = 'gbtree',learning_rate = 0.1),param_distributions = param
9 model.fit(np.array(train_avg2v),Y_train)#fitting the model
10
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n_jobs=1)]: Done 50 out of 50 | elapsed: 46.9min finished



```
Out[42]: RandomizedSearchCV(cv=TimeSeriesSplit(max_train_size=None, n_splits=5),
error_score='raise-deprecating',
estimator=XGBClassifier(base_score=0.5, booster='gbtree',
                        colsample_bylevel=1,
                        colsample_bynode=1,
                        colsample_bytree=1, gamma=0,
                        learning_rate=0.1, max_delta_step=0,
                        max_depth=3, min_child_weight=1,
                        missing=None, n_estimators=100,
                        n_jobs=1, nthread=None,
                        objective='binary:logistic',
                        random_state=0, reg_alpha=0,
                        reg_lambda=1, scale_pos_weight=1,
                        seed=None, silent=None, subsample=1,
                        verbosity=1),
iid='warn', n_iter=10, n_jobs=None,
param_distributions={'max_depth': [3, 9, 11, 13, 15],
                    'n_estimators': [10, 50, 100, 200,
                                     450]}},
pre_dispatch='2*n_jobs', random_state=None, refit=True,
return_train_score=True, scoring=None, verbose=1)
```

```
In [0]: 1 savetofile(model, 'model_avg2v_gb')
```

```
In [0]: 1 model_avgw2v_gb = openfromfile('model_avgw2v_gb')
```

```
In [45]: 1 plots(model_avgw2v_gb)
```

Best Hyperparameters are: {'n_estimators': 50, 'max_depth': 3}




```
In [0]: 1 best_depth_avgw2v_gb = model_avgw2v_gb.best_params_['max_depth']
        2 best_n_avgw2v_gb = model_avgw2v_gb.best_params_['n_estimators']
```

In [48]:

```
1  '''AUC ON TEST DATA'''
2  train_auc_avgw2v_gb,test_auc_avgw2v_gb,train_pred_proba,test_pred_proba,train_pred,test_pred = auc(best_dept
3
4  '''PLOTING THE ROC CURVE'''
5  curve(train_pred_proba,test_pred_proba)
6
7  '''Precision,recall and f1 score'''
8  metrics(test_pred)
9
10
11  '''Plotting the confusion matrix'''
12  print('Confusin matrix on train data')
13  plot_confusion_matrix(Y_train,train_pred)
14  print('Confusion matrix for test data')
15  plot_confusion_matrix(Y_test,test_pred)
```

AUC on train data is: 0.5544450279507798

AUC on test data is: 0.5021911936221565




scores on test data are:

precison score is 0.8421666666666666


recall_score is 1.0

f1 score is 0.9143219035555958

Confusin matrix on train data



Confusion matrix for test data



[6.4] Applying XGBOOST on TFIDF W2V, SET 4

```
In [32]: 1 train_tfidfw2v = openfromfile('computed/train_tfidfw2v')
          2 test_tfidfw2v = openfromfile('computed/test_tfidfw2v')
```

```
In [33]: 1 from sklearn.model_selection import RandomizedSearchCV
2 from xgboost import XGBClassifier
3
4 estimators = [10,50,100,200,450]#list of estimators that will be tuned
5 depths = [3,9,11,13,15]#tuning depth to avoid overfitting and underfitting
6 params = {'max_depth':depths,'n_estimators':estimators}#for passing as argument
7
8 tscv = TimeSeriesSplit(n_splits = 5)
9 model = RandomizedSearchCV(XGBClassifier(booster = 'gbtree',learning_rate = 0.1),param_distributions = param
10 model.fit(np.array(train_tfidf2v),Y_train)#fitting the model
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[Parallel(n_jobs=1)]: Done 50 out of 50 | elapsed: 128.9min finished


```
Out[33]: RandomizedSearchCV(cv=TimeSeriesSplit(n_splits=5), error_score='raise',
    estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
    max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
    n_jobs=1, nthread=None, objective='binary:logistic', random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=True, subsample=1),
    fit_params={}, iid=True, n_iter=10, n_jobs=1,
    param_distributions={'max_depth': [3, 9, 11, 13, 15], 'n_estimators': [10, 50, 100, 200, 450]},
    pre_dispatch='2*n_jobs', random_state=None, refit=True,
    return_train_score=True, scoring=None, verbose=1)
```

```
In [34]: 1 savetofile(model,'model_tfidf2v_gb')
```

```
In [35]: 1 model_tfidf2v_gb = openfromfile('model_tfidf2v_gb')
```


In [39]: 1 plots(model_tfidf2v_gb)

Best Hyperparameters are: {'n_estimators': 10, 'max_depth': 3}




In [55]: 1 best_depth_tfidf2v_gb = model_tfidf2v_gb.best_params_['max_depth']
 2 best_n_tfidf2v_gb = model_tfidf2v_gb.best_params_['n_estimators']

```
In [46]: 1 '''AUC ON TEST DATA'''
2 train_auc_tfidfw2v_gb,test_auc_tfidfw2v_gb,train_pred_proba,test_pred_proba,train_pred,test_pred = auc(best_
3     best_n_tfidfw2v_gb,np.array(train_tfidfw2v),test_tfidfw2v)
4
5 '''PLOTING THE ROC CURVE'''
6 curve(train_pred_proba,test_pred_proba)
7
8 '''Precision,recall and f1 score'''
9 metrics(test_pred)
10
11
12 '''Plotting the confusion matrix'''
13 print('Confusin matrix on train data')
14 plot_confusion_matrix(Y_train,train_pred)
15 print('Confusion matrix for test data')
16 plot_confusion_matrix(Y_test,test_pred)
```

AUC on train data is: 0.512827498261

AUC on test data is: 0.498112157921




scores on test data are:

precison score is 0.8412014299555645


recall_score is 0.9977016959898558

f1 score is 0.9127920677216452

Confusin matrix on train data



Confusion matrix for test data



[7] Conclusions

In [85]:

```

1  #create table using prettytable
2  from prettytable import PrettyTable
3
4  #table for random forest
5  table_rf = PrettyTable()
6  no = [1,2,3,4]
7  vectorizers = ['Bag of vectors', 'TFIDF', 'Average Word 2 vector', 'TFIDF Word 2 vector'] #all vectorizers
8  depths = ['50', '50', '15', '50'] #best parameters
9  estimators = ['450', '250', '250', '450']
10 AUC = ['0.501', '0.495', '0.903', '0.883'] #their respective auc scores
11
12 table_rf.add_column("SNo",no)
13 table_rf.add_column('Vectorizers',vectorizers)
14 table_rf.add_column('max depth',depths)
15 table_rf.add_column('no. of estimators',estimators)
16 table_rf.add_column('AUC on test',AUC)
17 print('\t\t\t\t\t Table for Random Forest')
18 print(table_rf)
19 print('\n\n\n')
20
21
22 table_gb = PrettyTable()
23
24 no_gb = [1,2,3,4]
25 vectorizers_gb = ['Bag of vectors', 'TFIDF', 'Average Word 2 vector', 'TFIDF Word 2 vector'] #all vectorizers
26 depths_gb = ['3', '3', '3', '3']
27 estimators_gb = ['100', '200', '50', '10']
28 AUC_gb = ['0.5008', '0.4984', '0.5021', '0.4981']
29
30 table_gb.add_column("SNo",no_gb)
31 table_gb.add_column('Vectorizers',vectorizers_gb)
32 table_gb.add_column('max depth',depths_gb)
33 table_gb.add_column('no. of estimators',estimators_gb)
34 table_gb.add_column('AUC on test for GBDT',AUC_gb)
35 print('\t\t\t\t\t Table for GBDT using XGBoost')
36 print(table_gb)
37
38
39

```

Table for Random Forest

```

+-----+-----+-----+-----+-----+

```

SNo	Vectorizers	max depth	no. of estimators	AUC on test
1	Bag of vectors	50	450	0.501
2	TFIDF	50	250	0.495
3	Average Word 2 vector	15	250	0.903
4	TFIDF Word 2 vector	50	450	0.883

Table for GBDT using XGBoost

SNo	Vectorizers	max depth	no. of estimators	AUC on test for GBDT
1	Bag of vectors	3	100	0.5008
2	TFIDF	3	200	0.4984
3	Average Word 2 vector	3	50	0.5021
4	TFIDF Word 2 vector	3	10	0.4981

In []:

1